

Introduction to Hydra in the Lightning-Hydra-Template

What is Hydra?

- Hydra is a **configuration framework** developed by **Facebook Research**.
- Used in the **Lightning-Hydra-Template** to:
 - Manage **hierarchical configurations**.
 - Modify settings dynamically via YAML or CLI.
- Hydra dynamically instantiates target classes (via the `_target_` key) in `train.py` using config files.

Configuration Structure

configs/

├── common # Folder with many configs (hydra, trainer, and paths in config_base.yaml)

├── data/ # Datasets

├── experiment/ # Experiments

├── model/ # Models

— train.yaml # Main training config

— eval.yaml # Main evaluation config

The main file is `configs/train.yaml`, which defines the default training configuration.

`train.yaml:5-10`

```
# @package _global_  
  
# specify here default configuration  
# order of defaults determines the order in which configs override each other  
defaults:  
  - _self_  
  - common/config_base  
  - data: mnist  
  - model: mnist  
  - /common/callbacks@callbacks: default # Notation explicite  
  - /common/logger@logger: mlflow # set logger here or use command line (e.g. `python train.py logger=tensorboard`)  
  
# experiment configs allow for version control of specific hyperparameters  
# e.g. best hyperparameters for given model and datamodule  
  - experiment: null  
  
...
```

We will load `configs/common/config_base.yaml`, which contains the trainer parameters and paths.

```
# ----- PATHS -----  
paths:  
  root_dir: ${oc.env:PROJECT_ROOT} # Root directory of the project, set via environment variable  
  data_dir: ${paths.root_dir}/data/ # Directory where the datasets are stored  
  log_dir: ${paths.root_dir}/logs/ # Directory where logs will be saved  
  output_dir: ${hydra:runtime.output_dir} # Output directory generated by Hydra for each run  
  work_dir: ${hydra:runtime.cwd} # Current working directory where the job is launched  
  
# ----- TRAINER -----  
trainer:  
  _target_: lightning.pytorch.trainer.Trainer #target class for trainer  
  default_root_dir: ${paths.output_dir}  
  min_epochs: 1  
  max_epochs: 10  
  accelerator: gpu  
  ...
```

We also load the datamodule (in this case, MNIST): `configs/data/mnist.yaml`

```
_target_: src.data.mnist_datamodule.MNISTDataModule #target class for datamodule
data_dir: ${paths.data_dir}
batch_size: 128 # Needs to be divisible by the number of devices (e.g., if in a distributed setup)
train_val_test_split: [55_000, 5_000, 10_000]
num_workers: 0
pin_memory: False
```

We load our model (MNIST): `configs/model/mnist.yaml`

```
_target_: src.models.mnist_module.MNISTLitModule #target class for the model

optimizer:
  _target_: torch.optim.Adam #target class for the optimizer
  _partial_: true
  lr: 0.001
  weight_decay: 0.0

scheduler:
  _target_: torch.optim.lr_scheduler.ReduceLROnPlateau #target class for the scheduler
  _partial_: true
  mode: min
  factor: 0.1
  patience: 10

net:
  _target_: src.models.components.simple_dense_net.SimpleDenseNet #target class selfmade for the net
  input_size: 784
  lin1_size: 64
  lin2_size: 128
  lin3_size: 64
  output_size: 10

# compile model for faster training with pytorch 2.0
compile: false

# learning rate scheduler update interval: "epoch" or "step"
lr_interval: epoch
```

We also load the callbacks: `configs/common/callbacks/default.yaml`

```
defaults:
  - model_checkpoint
  - early_stopping
  - model_summary
  - rich_progress_bar
  - callback_images
  - _self_

...
```

This file lists the callbacks we want to use by default.

Finally, we load the logger from `configs/common/logger` — here are the available options:

✓ logger

- ! `many_loggers.yaml`

- ! `mlflow.yaml`

- ! `tensorboard.yaml`

⚙️ How to Modify Configuration Variables

You can change values in several ways:

1. Direct Modification in YAML Files

Change values directly in YAML files. For example, modify the number of max epochs in `configs/common/config_base.yaml`:

```
trainer:  
  _target_: lightning.pytorch.trainer.Trainer  
  default_root_dir: ${paths.output_dir}  
  min_epochs: 1  
  max_epochs: 10  
  accelerator: gpu  
  devices: 1  
  check_val_every_n_epoch: 1  
  deterministic: false
```

2. Override via Command Line

Change one or more parameters without modifying YAML files:

```
python src/train.py model.optimizer.lr=0.001 data.batch_size=64
```

3. Alternative Configurations

Use a different model or logger:

```
python src/train.py model=autre_modele common/logger@logger=nom_logger
```

IMPORTANT: To override logger or callbacks, use the following syntax:

```
python src/train.py common/logger@logger=nom_logger
```

```
python src/train.py common/callbacks@callbacks=nom_callbacks
```

Experiment Configurations

Goal:

Define a full experiment configuration in a single YAML file
(to avoid manually changing all individual config files):

- Model
- Dataset
- Callbacks
- Number of epochs
- Seed

Example: configs/experiment/mon_experience.yaml

```
# @package _global_
defaults:
  - override /data: mnist
  - override /model: mnist
  - override /common/callbacks@callbacks: default
  - override /common/logger@logger: mlflow

# Paramètres spécifiques à cette expérience
tags: ["mnist"]

# Surcharger des paramètres spécifiques
trainer:
  max_epochs: 10
  min_epochs: 5
  gradient_clip_val: 0.5
  accelerator: cpu
  devices: 1

model:
  optimizer:
    lr: 0.002
  net:
    lin1_size: 128
    lin2_size: 256
    lin3_size: 64

data:
  batch_size: 64
```

Then launch your training with:

```
python src/train.py experiment=mon_experience
```

This will load:

- MNIST dataset
- MNIST model
- Default callbacks
- Seed 12345
- 20 epochs
- Learning rate of 0.002

So, before overriding parameters in an experiment, refer to the original config file to check the variable structure.

Final Check

If everything went well...

You should see this image generated at the start of training in the terminal:

```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
python3 + v [ ] [ ] ... ^ x

TPU available: False, using: 0 TPU cores
HPU available: False, using: 0 HPUs
[2025-05-19 11:13:41,532][__main__][INFO] - [rank: 0] Logging hyperparameters!
[2025-05-19 11:13:41,598][__main__][INFO] - [rank: 0] Starting training!



|    | Name        | Type             | Params | Mode  |
|----|-------------|------------------|--------|-------|
| 0  | net         | SimpleDenseNet   | 151 K  | train |
| 1  | net.model   | Sequential       | 151 K  | train |
| 2  | net.model.0 | Linear           | 100 K  | train |
| 3  | net.model.1 | BatchNorm1d      | 256    | train |
| 4  | net.model.2 | ReLU             | 0      | train |
| 5  | net.model.3 | Linear           | 33.0 K | train |
| 6  | net.model.4 | BatchNorm1d      | 512    | train |
| 7  | net.model.5 | ReLU             | 0      | train |
| 8  | net.model.6 | Linear           | 16.4 K | train |
| 9  | net.model.7 | BatchNorm1d      | 128    | train |
| 10 | net.model.8 | ReLU             | 0      | train |
| 11 | net.model.9 | Linear           | 650    | train |
| 12 | criterion   | CrossEntropyLoss | 0      | train |
| 13 | train_loss  | MeanMetric       | 0      | train |
| 14 | val_loss    | MeanMetric       | 0      | train |
| 15 | test_loss   | MeanMetric       | 0      | train |



Trainable params: 151 K
Non-trainable params: 0
Total params: 151 K
Total estimated model params size (MB): 0
Modules in train mode: 16
Modules in eval mode: 0
/home/guillaume/.local/lib/python3.10/site-packages/lightning/pytorch/trainer/connectors/data_connector.py:425: The 'val_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the 'num_workers' argument to 'num_workers=7' in the 'DataLoader' to improve performance.
/home/guillaume/.local/lib/python3.10/site-packages/lightning/pytorch/trainer/connectors/data_connector.py:425: The 'train_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the 'num_workers' argument to 'num_workers=7' in the 'DataLoader' to improve performance.
Epoch 2/9 537/860 0:00:11 • 0:00:07 47.55it/s v_num: c816 val/loss: 0.089 train/loss: 0.101

```